



# Software Defined Networking

Dr. Nick Feamster  
Professor

---

*In this course, you will learn about software defined networking and how it is changing the way communications networks are managed, maintained, and secured.*

## **This Module: Programmable Data Plane**

- ◎ Two Lessons
  - Programming the data plane: Click
  - **Scaling programmable data planes**
    - **Making software faster**
    - Making hardware more programmable
- ◎ **Optional** programming assignment (in Click)
- ◎ Quiz on Concepts

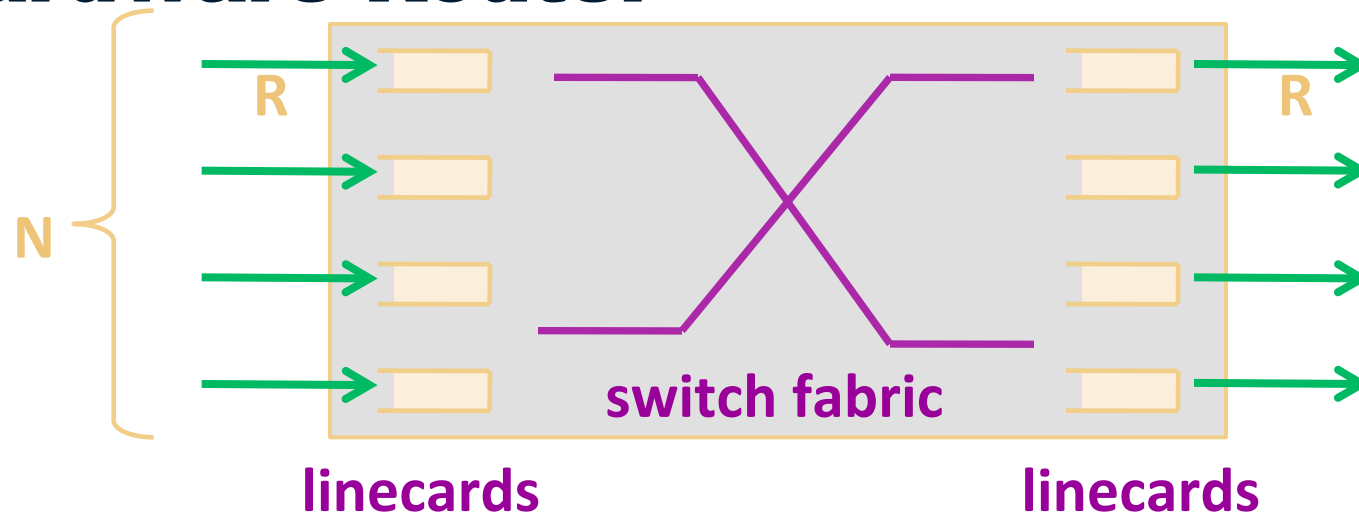
## Motivation

- ⦿ Many new protocols require data-plane changes.
  - Examples: OpenFlow, Path Splicing, AIP, ...
- ⦿ Protocols must forward packets at acceptable speeds.
- ⦿ May need to run in parallel with existing protocols
- ⦿ **Need:** Platform for developing new network protocols that
  - Forwards packets at high speed
  - Runs multiple data-plane protocols in parallel

## Existing Approaches


- ◎ Develop **custom software**
  - **Advantage:** Flexible, easy to program
  - **Disadvantage:** Slow forwarding speeds
- ◎ Develop modules in **custom hardware**
  - **Advantage:** Excellent performance
  - **Disadvantage:** Long development cycles, rigid
- ◎ Develop in **programmable hardware**
  - **Advantage:** Flexible and fast
  - **Disadvantage:** Programming is difficult

## Hardware Router



- Processing at rate  $\sim R$  per line card
- Switching at rate  $N \times R$  by switch fabric

## RouteBricks: Linecards on Servers



- Processing at rate  $\sim R$  per server
- Switching at rate  $\sim R$  per server

## Requirements




- Internal link rates  $< R$
- Per-server processing rate:  $c \times R$
- Per-server fanout: constant

## Challenges


- ◎ **Limited internal link rates:** Internal links can't exceed external link rates
- ◎ **Limited per-node processing rate:** Desire to use commodity hardware
- ◎ **Limited per-node fanout:** Due to limited NIC slots/ports



## Strawman Approach




## Strawman Approach




- ⦿  $N$  external links of capacity  $R$
- ⦿  $N^2$  internal links of capacity  $R$

## Valiant Load Balancing



- Per-server processing rate:  $3R$
- With uniform traffic (avoiding first phase):  $2R$

## Each Server Must Also Be Fast




◎ First try: 1.3 Gbps

## Problem #1: Bookkeeping


- ⦿ **Managing packet descriptors**
  - moving between NIC and memory
  - updating descriptor rings
- ⦿ **Solution: batch packet operations**
  - NIC batches multiple packet descriptors
  - CPU polls for multiple packets
  - Cost: increased latency

## Single-Server Performance




- First try: 1.3 Gbps
- With batching: **3 Gbps**

## Problem #2: Queue Access



- ⦿ Rule #1: 1 core per queue (avoids locking)
- ⦿ Rule #2: 1 core per packet (faster)

## Single-Server Performance



- First try: 1.3 Gbps
- With batching: 3 Gbps
- With multiple queues: **9.7 Gbps**



## Fast Software Forwarding: Other Tricks

- ⦿ Large packet buffers to hold multiple packets
- ⦿ Batch processing
- ⦿ Ethernet GRE (to avoid complicated lookup)
- ⦿ Avoiding lookups on bridge between virtual interfaces and physical interfaces

Han, Sangjin, et al. "PacketShader: a GPU-accelerated software router." *ACM SIGCOMM Computer Communication Review* 40.4 (2010): 195-206.

Bhatia, Sapan, et al. "Trellis: A platform for building flexible, fast virtual networks on commodity hardware." *Proceedings of the 2008 ACM CoNEXT Conference*. ACM, 2008.

## Summary

- ◎ **Scalability: Make the software faster**
  - **Software routers can be fast!**
- ◎ General purpose infrastructure is capable of fast forwarding performance
  - The low-level details, optimizations matter
- ◎ Other efforts underway
  - Intel DPDK